

TEAM SANKALPA

Hybrid Local-Fireworks Routing Agent

A highly competitive, token-efficient routing agent designed to balance maximum accuracy with minimal API costs - intelligently matching every prompt to the right level of intelligence.

Made with GAMMA

The Cost vs. Accuracy Dilemma



Large Models

Powerful but expensive and slow, overkill for simple tasks like sentiment analysis or factual QA.

Small Models

Cheap and fast, but struggle with complex reasoning, multi-step logic, and nuanced code generation.

The Goal

Achieve maximum accuracy without paying a premium for every single prompt.

Made with GAMMA

Multi-Tier Cost-Aware Routing

Intelligently matching each task to the right level of intelligence — no more, no less.

Tier 1 : Baseline

Handles basic language tasks efficiently: sentiment analysis, NER, factual QA, and summarization.

Tier 2 : Intermediate

Balances cost and capability for multi-step reasoning tasks and standard code generation.

Tier 3 : Premium

Reserved for complex tasks where hallucinations are costly: intricate logic, math, and code debugging.

Made with GAMMA

Advanced Optimizations

Pushing the limits of efficiency at every layer of the system.



Verifier Safety Net

Fast heuristic checks discard hallucinated outputs, automatically falling back to premium tiers to protect accuracy.



Token Optimization

Aggressively minimizes costs using strict output caps and concise formatting requirements.



TLE Panic Mode

Instantly routes to premium models when the evaluation time budget drops dangerously low, preventing timeouts.



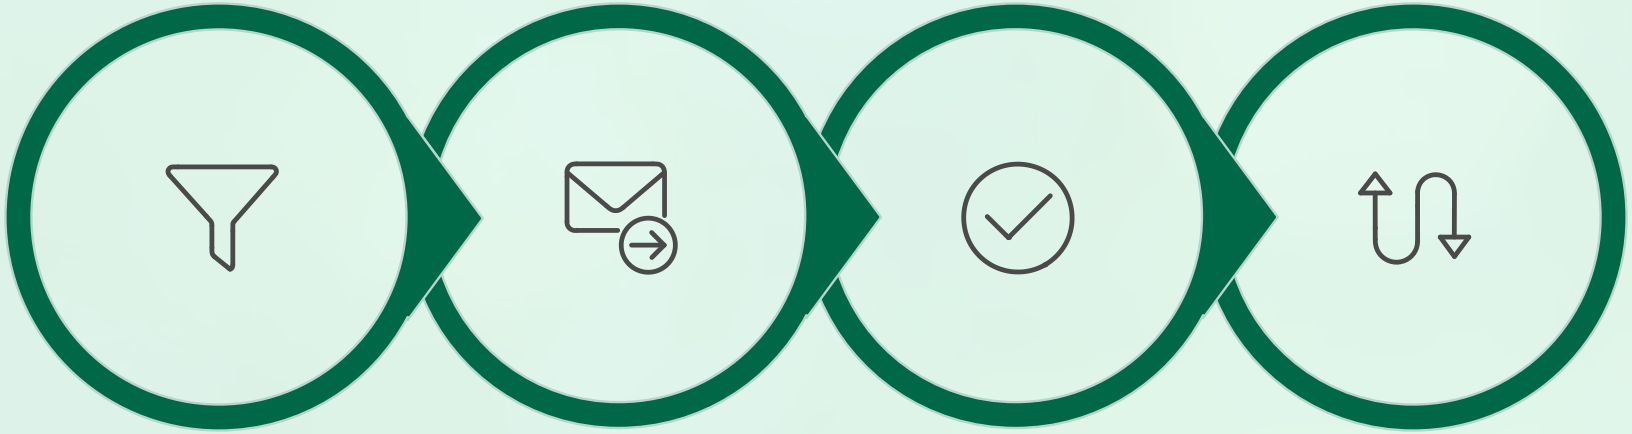
Dynamic Model Handling

Uses Regex heuristics to infer model capabilities and handles unknown inputs gracefully without crashing.

Made with GAMMA

The Routing Pipeline

How a prompt moves through the system — from classification to confident output.



Classification

Attempt

Validation

Escalation

Each prompt is classified by intent, routed to the lowest-cost tier capable of handling it, validated for correctness, and escalated only when necessary — keeping costs minimal without sacrificing reliability.

Made with GAMMA



RESILIENCE

Built for Strict, Unpredictable Environments

Memory Constrained

Engineered to operate flawlessly in restrictive 4GB RAM environments.

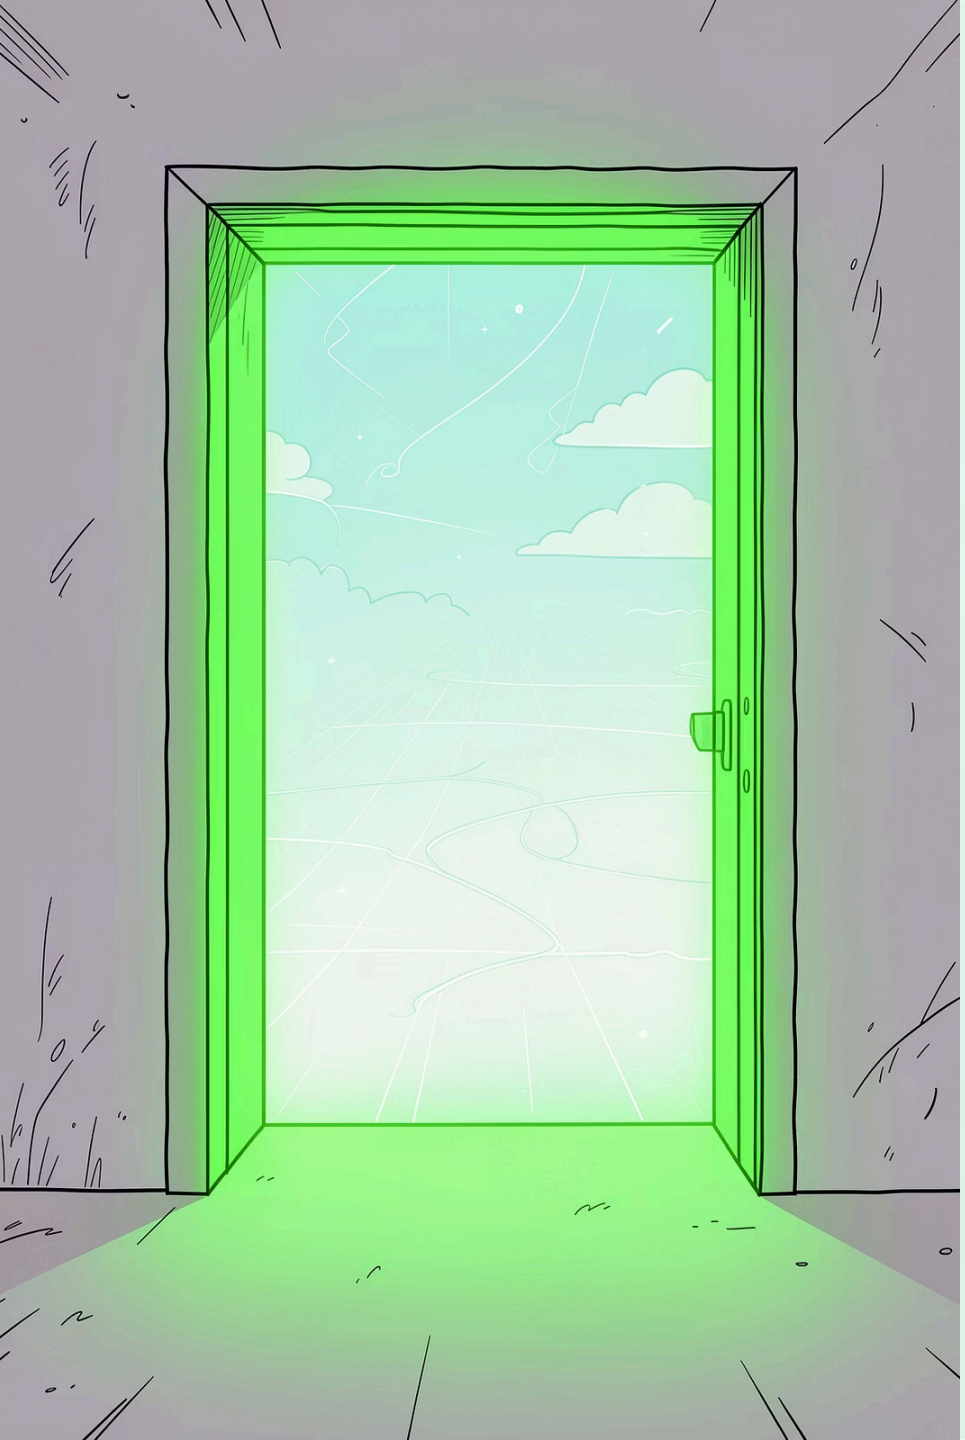
Fault Tolerant

Gracefully catches network 404s, API timeouts, and malformed JSON — the container never exits with a fatal error.

Isolated Execution

Fully containerized with no external dependencies required at runtime.

Made with GAMMA



Ready to Route?

Experience the optimal balance of intelligence and efficiency. The Hybrid Local-Fireworks Routing Agent is open source and ready to integrate.

🔗 github.com/Amanjha112113/hybrid-llm-router

Token-Efficient

Route smartly, spend less.

Accuracy-First

Never compromise on quality.

Production-Ready

Containerized and resilient.

Made with GAMMA